

Name:Viviyan Richards W

Roll no:205229133

Lab4. Pandas Grouping and Aggregation

Import necessary modules

In [1]:

```
import pandas as pd
```

In [2]:

```
data = pd.read_csv("thanksgiving-2015-poll-data.csv", encoding="Latin-1")
```


In [3]:

```
data.head()
```

Out[3]:

	RespondentID	Do you celebrate Thanksgiving?	What is typically the main dish at your Thanksgiving dinner?	What is typically the main dish at your Thanksgiving dinner? - Other (please specify)	How is the main dish typically cooked?	How is the main dish typically cooked? - Other (please specify)	What kind of stuffing/dressing do you typically have?
0	4337954960	Yes	Turkey	NaN	Baked	NaN	Bread-based
1	4337951949	Yes	Turkey	NaN	Baked	NaN	Bread-based
2	4337935621	Yes	Turkey	NaN	Roasted	NaN	Rice-based
3	4337933040	Yes	Turkey	NaN	Baked	NaN	Bread-based
4	4337931983	Yes	Tofurkey	NaN	Baked	NaN	Bread-based

5 rows × 65 columns



In [4]:

```
data.shape
```

Out[4]:

```
(1058, 65)
```

What are unique values of "Do you celebrate Thanksgiving?" column?

```
In [5]: data["Do you celebrate Thanksgiving?"].unique()
```

```
Out[5]: array(['Yes', 'No'], dtype=object)
```

View all column names (top 5)

```
In [6]: data.columns[:5]
```

```
Out[6]: Index(['RespondentID', 'Do you celebrate Thanksgiving?',  
              'What is typically the main dish at your Thanksgiving dinner?',  
              'What is typically the main dish at your Thanksgiving dinner? - Other (p  
              lease specify)',  
              'How is the main dish typically cooked?'],  
             dtype='object')
```

Applying functions to Sreies

How many male, female and NaN in "What is your gender?" column

```
In [7]: data["What is your gender?"].unique()  
data["What is your gender?"].value_counts(dropna=False)
```

```
Out[7]: Female      544  
Male        481  
NaN          33  
Name: What is your gender?, dtype: int64
```

Let apply a user defined function to each value in the What is your gender? column to transform Male to 0 and female to 1

```
In [8]: import math  
  
def gender_code(gender_val):  
    if isinstance(gender_val, float) and math.isnan(gender_val):  
        return gender_val  
    return int(gender_val.lower().strip() == "female")
```

Apply gender_code() to What is your gender? column

```
In [9]: gender_codes = data["What is your gender?"].apply(gender_code)
print(type(gender_codes), gender_codes.head())
```

```
<class 'pandas.core.series.Series'> 0    0.0
1    1.0
2    0.0
3    0.0
4    0.0
Name: What is your gender?, dtype: float64
```

Now, count male and females as 0s and 1s. How many in "gender" column?

```
In [10]: data["gender"] = gender_codes
data["gender"].value_counts(dropna=False)
```

```
Out[10]: 1.0    544
0.0    481
NaN      33
Name: gender, dtype: int64
```

Applying functions to DataFrames

Check the data type of each column in data using a lambda function. Just visualize data types of first 5 columns

```
In [11]: def get_type(row):
return row.dtype
data.apply(get_type).head()
```

```
Out[11]: RespondentID
int64
Do you celebrate Thanksgiving?
object
What is typically the main dish at your Thanksgiving dinner?
object
What is typically the main dish at your Thanksgiving dinner? - Other (please specify)
object
How is the main dish typically cooked?
object
dtype: object
```

DATA CLEANING - Let us clean up Income column

We need to convert string values representing income in "How much total combined money did all members of your HOUSEHOLD earn last year" column into numeric values. Check the unique values first

```
In [12]: column_name = "How much total combined money did all members of your HOUSEHOLD ea
data[column_name].value_counts(dropna=False)
```

```
Out[12]: $25,000 to $49,999      180
        Prefer not to answer    136
        $50,000 to $74,999      135
        $75,000 to $99,999      133
        $100,000 to $124,999    111
        $200,000 and up         80
        $10,000 to $24,999      68
        $0 to $9,999           66
        $125,000 to $149,999    49
        $150,000 to $174,999    40
        NaN                    33
        $175,000 to $199,999    27
        Name: How much total combined money did all members of your HOUSEHOLD earn last
        year?, dtype: int64
```

```
In [13]: import numpy as np

def clean_income(value):
    if value == "$200,000 and up":
        return 200000
    elif value == "Prefer not to answer":
        return np.nan
    elif isinstance(value, float) and math.isnan(value):
        return np.nan
    value = value.replace(",", "").replace("$", "")
    income_low, income_high = value.split(" to ")

    return (int(income_high) + int(income_low)) / 2
```

Now apply this function to the "How much total combined money did all members of your HOUSEHOLD earn last year?" column and put it in new column "income"

```
In [14]: column_name = "How much total combined money did all members of your HOUSEHOLD ea
data["income"] = data[column_name].apply(clean_income)
data["income"].head()
```

```
Out[14]: 0      87499.5
        1      62499.5
        2       4999.5
        3     200000.0
        4     112499.5
        Name: income, dtype: float64
```

Grouping Data with Pandas

Check unique values in column, "What type of cranberry saucedo you typically have?" first.

```
In [15]: data["What type of cranberry saucedo you typically have?"].value_counts()
```

```
Out[15]: Canned          502
Homemade        301
None            146
Other (please specify)  25
Name: What type of cranberry saucedo you typically have?, dtype: int64
```

Create a datafrme by filtering values "Homemade"

```
In [16]: column_name = "What type of cranberry saucedo you typically have?"
homemade_mask = data[column_name] == "Homemade"
print(homemade_mask.head())

homemade = data[homemade_mask]
```

```
0    False
1    False
2     True
3     True
4    False
Name: What type of cranberry saucedo you typically have?, dtype: bool
```

Create another datafrme by filtering values "Canned"

```
In [17]: canned = data[data[column_name] == "Canned"]
```

Now print mean income of homemade_df and canned_df for these two groups of people

```
In [18]: print(homemade["income"].mean())
print(canned["income"].mean())
```

```
94878.1072874494
83823.40340909091
```

Use groupby() and aggregation() to find out "Who earns more income?"

Split dataset based on "What type of cranberry saucedo you typically have?" column automatically into groups based on unique values

```
In [19]: column_name = "What type of cranberry saucedo you typically have?"
grouped = data.groupby(column_name)
grouped
```

```
Out[19]: <pandas.core.groupby.generic.DataFrameGroupBy object at 0x0000026CF97F2AC0>
```

List out all groups that are created by groupby()

In [20]: `grouped.groups`

Out[20]: {'Canned': [4, 6, 8, 11, 12, 15, 18, 19, 26, 27, 38, 43, 48, 53, 58, 59, 60, 68, 69, 71, 74, 76, 79, 80, 86, 87, 89, 90, 91, 97, 103, 106, 107, 109, 115, 116, 118, 119, 123, 127, 129, 130, 132, 135, 136, 137, 140, 141, 143, 144, 145, 150, 153, 155, 156, 157, 158, 159, 161, 162, 163, 166, 167, 168, 169, 173, 179, 180, 181, 182, 184, 186, 190, 192, 193, 195, 198, 199, 200, 204, 205, 207, 209, 210, 211, 212, 213, 215, 217, 218, 220, 222, 224, 226, 229, 230, 231, 239, 243, 245, ...], 'Homemade': [2, 3, 5, 7, 13, 14, 16, 20, 21, 23, 25, 28, 30, 32, 33, 37, 39, 42, 44, 46, 52, 54, 56, 57, 62, 64, 66, 70, 82, 83, 85, 88, 93, 94, 96, 98, 101, 102, 108, 110, 111, 112, 114, 120, 122, 128, 134, 138, 139, 152, 165, 171, 172, 174, 175, 176, 177, 178, 183, 188, 189, 194, 201, 202, 203, 208, 219, 223, 225, 232, 234, 235, 236, 238, 241, 242, 244, 246, 248, 254, 255, 256, 259, 261, 262, 263, 264, 268, 281, 285, 286, 287, 290, 291, 292, 295, 298, 300, 302, 303, ...], 'None': [0, 17, 24, 29, 34, 36, 40, 47, 49, 51, 55, 61, 67, 72, 73, 77, 78, 81, 92, 99, 100, 104, 105, 117, 121, 124, 126, 131, 133, 142, 146, 148, 149, 160, 164, 185, 187, 191, 197, 227, 228, 237, 240, 274, 275, 319, 321, 329, 337, 362, 370, 377, 391, 395, 406, 409, 414, 417, 421, 437, 439, 466, 480, 491, 492, 495, 505, 514, 526, 529, 532, 537, 540, 553, 560, 564, 571, 573, 580, 584, 591, 594, 598, 602, 605, 606, 609, 610, 618, 626, 631, 639, 647, 658, 672, 673, 684, 700, 701, 716, ...], 'Other (please specify)': [1, 9, 154, 216, 221, 233, 249, 265, 301, 336, 380, 435, 444, 447, 513, 550, 749, 750, 784, 807, 860, 872, 905, 1000, 1007]}

In [21]: `grouped.size()`

Out[21]: What type of cranberry saucedo you typically have?
 Canned 502
 Homemade 301
 None 146
 Other (please specify) 25
 dtype: int64

In [22]: `for name, group in grouped:
 print(name)
 print('\t', group.shape)
 print('\t', type(group))`

Canned
 (502, 67)
 <class 'pandas.core.frame.DataFrame'>
 Homemade
 (301, 67)
 <class 'pandas.core.frame.DataFrame'>
 None
 (146, 67)
 <class 'pandas.core.frame.DataFrame'>
 Other (please specify)
 (25, 67)
 <class 'pandas.core.frame.DataFrame'>

```
In [23]: grouped["income"]
```

```
Out[23]: <pandas.core.groupby.generic.SeriesGroupBy object at 0x0000026CF770A940>
```

```
In [24]: grouped["income"].size()
```

```
Out[24]: What type of cranberry saucedo you typically have?
Canned          502
Homemade        301
None            146
Other (please specify)  25
Name: income, dtype: int64
```

Aggregating values in groups

Now, find out average income

```
In [25]: grouped["income"].agg(np.mean)
```

```
Out[25]: What type of cranberry saucedo you typically have?
Canned          83823.403409
Homemade        94878.107287
None            78886.084034
Other (please specify)  86629.978261
Name: income, dtype: float64
```

If you want to consider all numeric attributes and find the mean for each group for every column in data, you can do as below.

```
In [26]: grouped.agg(np.mean)
```

```
Out[26]:
```

	RespondentID	gender	income
What type of cranberry saucedo you typically have?			
Canned	4.336699e+09	0.552846	83823.403409
Homemade	4.336792e+09	0.533101	94878.107287
None	4.336765e+09	0.517483	78886.084034
Other (please specify)	4.336763e+09	0.640000	86629.978261

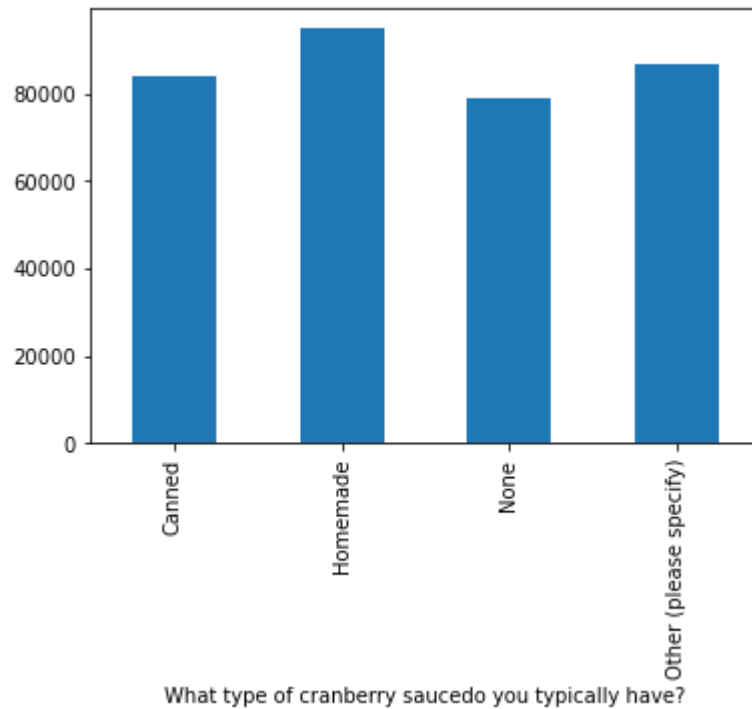
Plotting the results

What is the average income of each category?

```
In [27]: %matplotlib inline

sauce = grouped.agg(np.mean)
sauce["income"].plot(kind="bar")
```

Out[27]: <AxesSubplot:xlabel='What type of cranberry saucedo you typically have? '>



Aggregation with multiple columns

Find the average income of people who eat Homemade cranberry sauce and Tofurkey


```
In [28]: grouped = data.groupby(
    ["What type of cranberry saucedo you typically have?",
     "What is typically the main dish at your Thanksgiving dinner?"])
grouped.agg(np.mean)
```

Out[28]:

		RespondentID	gender	income
What type of cranberry saucedo you typically have?	What is typically the main dish at your Thanksgiving dinner?			
Canned	Chicken	4.336354e+09	0.333333	80999.600000
	Ham/Pork	4.336757e+09	0.642857	77499.535714
	I don't know	4.335987e+09	0.000000	4999.500000
	Other (please specify)	4.336682e+09	1.000000	53213.785714
	Roast beef	4.336254e+09	0.571429	25499.500000
	Tofurkey	4.337157e+09	0.714286	100713.857143
	Turkey	4.336705e+09	0.544444	85242.682045
	Chicken	4.336540e+09	0.750000	19999.500000
	Ham/Pork	4.337253e+09	0.250000	96874.625000
	I don't know	4.336084e+09	1.000000	NaN
Homemade	Other (please specify)	4.336863e+09	0.600000	55356.642857
	Roast beef	4.336174e+09	0.000000	33749.500000
	Tofurkey	4.336790e+09	0.666667	57916.166667
	Turducken	4.337475e+09	0.500000	200000.000000
	Turkey	4.336791e+09	0.531008	97690.147982
	Chicken	4.336151e+09	0.500000	11249.500000
	Ham/Pork	4.336680e+09	0.444444	61249.500000
	I don't know	4.336412e+09	0.500000	33749.500000
	Other (please specify)	4.336688e+09	0.600000	119106.678571
	Roast beef	4.337424e+09	0.000000	162499.500000
None	Tofurkey	4.336950e+09	0.500000	112499.500000
	Turducken	4.336739e+09	0.000000	NaN
	Turkey	4.336784e+09	0.523364	74606.275281
	Ham/Pork	4.336465e+09	1.000000	87499.500000
	Other (please specify)	4.337335e+09	0.000000	124999.666667
	Tofurkey	4.336122e+09	1.000000	37499.500000
	Turkey	4.336724e+09	0.700000	82916.194444
Other (please specify)				

Aggregation with multiple functions

Find sum, mean and standard deviation of each group in the income column of grouped dataframe

```
In [29]: grouped["income"].agg([np.mean, np.sum, np.std]).head()
```

Out[29]:

		mean	sum	std
What type of cranberry saucedo you typically have?	What is typically the main dish at your Thanksgiving dinner?			
	Chicken	80999.600000	404998.0	75779.481062
	Ham/Pork	77499.535714	1084993.5	56645.063944
	I don't know	4999.500000	4999.5	NaN
	Other (please specify)	53213.785714	372496.5	29780.946290
Canned	Roast beef	25499.500000	127497.5	24584.039538

Find the number of people who live in each area type (Rural, Suburban, etc) who eat different kinds of main dishes for Thanksgiving

```
In [30]: grouped = data.groupby("How would you describe where you live?")["What is typically the main dish at your Thanksgiving dinner?"]
grouped.apply(lambda x: x.value_counts())
```

```
Out[30]: How would you describe where you live?
Rural
          Turkey      189
          Other (please specify)    9
          Ham/Pork      7
          I don't know      3
          Tofurkey      3
          Turducken      2
          Chicken      2
          Roast beef      1
Suburban
          Turkey      449
          Ham/Pork      17
          Other (please specify)    13
          Tofurkey      9
          Roast beef      3
          Chicken      3
          I don't know      1
          Turducken      1
Urban
          Turkey      198
          Other (please specify)    13
          Tofurkey      8
          Chicken      7
          Roast beef      6
          Ham/Pork      4
Name: What is typically the main dish at your Thanksgiving dinner?, dtype: int64
```